

Comparison and Analysis of GAN Variants on CIFAR-10

Sr.No	Name	BITS ID
1.	Hari Shankar Jaiswal	2023aa05106
2.	Warun Kumar	2023aa05244
3.	Gandhi Disha	2023aa05388

1. Introduction

This report presents a detailed analysis of different Generative Adversarial Networks (GANs) trained on the CIFAR-10 dataset. The following GAN architectures were implemented and evaluated:

- **Task 1:** Conditional Wasserstein GAN (WGAN)
- **Task 2:** Spectrally Normalized GAN (SNGAN)
- **Task 3:** Self-Attention GAN (SAGAN) without spectral normalization and TTUB
- **Task 4:** Complete Self-Attention GAN (SAGAN)

Each model was evaluated based on the **Inception Score (IS)** and **Fréchet Inception Distance (FID)**, and their generated images were qualitatively assessed.

2. Comparative Table of GAN Variants

Model	Stability	Mode Collapse	Training Speed	IS Score	FID Score	Image Quality
Conditional WGAN	Moderate	Low	Slow	1.1573 ± 0.0065	63.8121	Good
SNGAN	High	Very Low	Moderate	1.1435 ± 0.0046	35.1275	Very Good
SAGAN (No Spectral Norm)	Low	High	Fast	1.1440 ± 0.0058	92.0718	Poor
Complete SAGAN	Very High	Very Low	Moderate	1.1035 ± 0.0039	85.2918	Good

3. Detailed Analysis of Each Model

3.1 Conditional WGAN (Task 1)

- **Architecture:** Uses a Wasserstein loss function with gradient penalty and conditions the generator and discriminator on class labels.
- **Observations:**
 - More stable than vanilla GANs but requires careful tuning.
 - Image quality is **good** but not as sharp as in later models.
 - Mode collapse is **low** due to the conditional setting.
 - Training is relatively **slow** due to the gradient penalty computation.
- **Metrics:** $IS = 1.1573 \pm 0.0065$, $FID = 63.8121$
- **Conclusion:** Performs well but is outperformed by SNGAN and SAGAN in terms of image quality and FID score.

3.2 SNGAN (Task 2)

- **Architecture:** Uses **spectral normalization** to stabilize training and improve discriminator regularization.
- **Observations:**
 - Provides **significantly better stability** than WGAN.
 - Mode collapse is almost **eliminated**.
 - Generates **sharper** and **more diverse** images.
 - Training is **faster** than WGAN.
- **Metrics:** $IS = 1.1435 \pm 0.0046$, $FID = 35.1275$
- **Conclusion:** SNGAN **outperforms WGAN** in stability, training efficiency, and image quality. Image quality in this SNGAN is really good.

3.3 SAGAN without Spectral Normalization and TTUB (Task 3)

- **Architecture:** SAGAN but without **spectral normalization** and **TTUB (Truncated Truncated Upper Bound)**.
- **Observations:**
 - **Unstable** training, leading to frequent mode collapse.
 - Poor quality and **noisy artifacts** in generated images.
 - Fast training but at the cost of low-quality outputs.
- **Metrics:** $IS = 1.1440 \pm 0.0058$, $FID = 92.0718$
- **Conclusion:** Removing spectral normalization negatively impacts performance, making training unstable.

3.4 Complete SAGAN (Task 4)

- **Architecture:** Incorporates **self-attention** along with **spectral normalization** to improve feature representation.
- **Observations:**
 - **Good image quality** with **sharp, high-resolution features**.
 - **Most stable training** with **minimal mode collapse**.
 - Generates **realistic images** with **high diversity**.
 - Also observed problem of Overfitting due to model collapse.
 - Training is **moderately fast** but provides superior results.

- **Metrics:** IS = 1.1035 ± 0.0039 , FID = 85.2918
- **Conclusion:** The complete SAGAN is the **best-performing** model in terms of both **quantitative scores** and **visual quality**.

4. Final Observations and Takeaways

4.1 Performance Summary

- **SAGAN (complete)** is the best model overall, but model collapses resulting into overfitting..
- **SNGAN** is a great balance between stability and efficiency.
- **Conditional WGAN** performs well but is outclassed by newer architectures.
- **SAGAN without spectral normalization** suffers from severe instability and poor results.

4.2 Lessons Learned

- **Spectral normalization** is crucial for stabilizing training.
- **Self-attention mechanisms** improve image quality by capturing long-range dependencies.
- **Gradient penalty (WGAN)** improves training but slows down convergence.

4.3 Recommendations for Future Improvements

- Use a **larger dataset** to improve generalization.
- Experiment with **hybrid models**, combining elements from different architectures.
- Apply **advanced regularization techniques** to enhance training stability.

5. Conclusion

This study demonstrates the **progression of GAN architectures** from traditional WGANs to modern self-attention GANs. The findings highlight the importance of **spectral normalization, attention mechanisms, and architectural design choices** in improving GAN performance. **SAGAN with self-attention and spectral normalization achieves the best balance of stability, diversity, and realism.**